



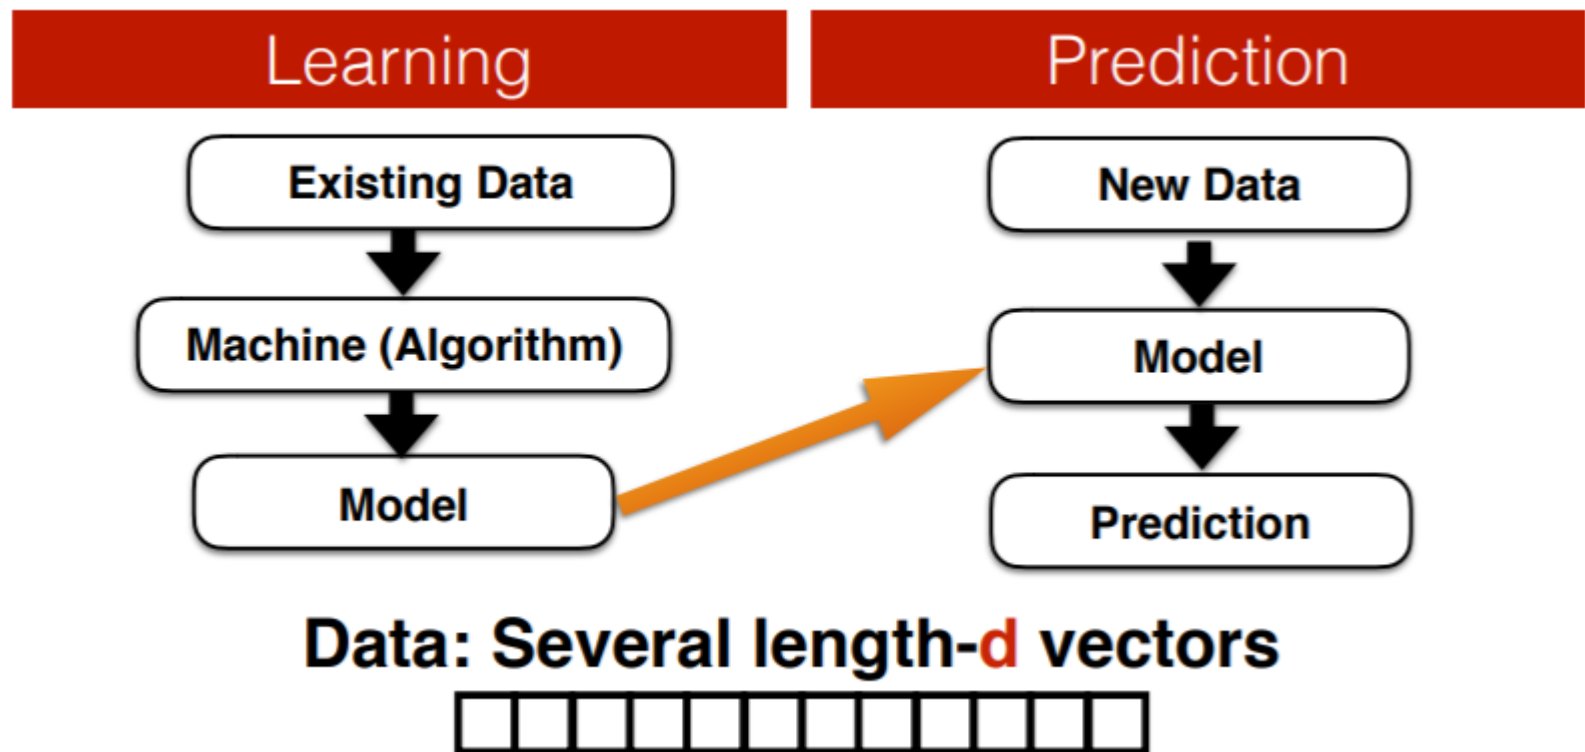
Feature Engineering and Preprocessing

shaileshsivan@gmail.com

Overview

- Dimensionality Reduction
- Components of Dimensionality Reduction
- Methods of Dimensionality Reduction
- Feature Reduction Iris Dataset
- Preprocessing

what machine learning does?



Some Facts

Expectation :

- We have good enough data
- Do focus on designing better algorithms

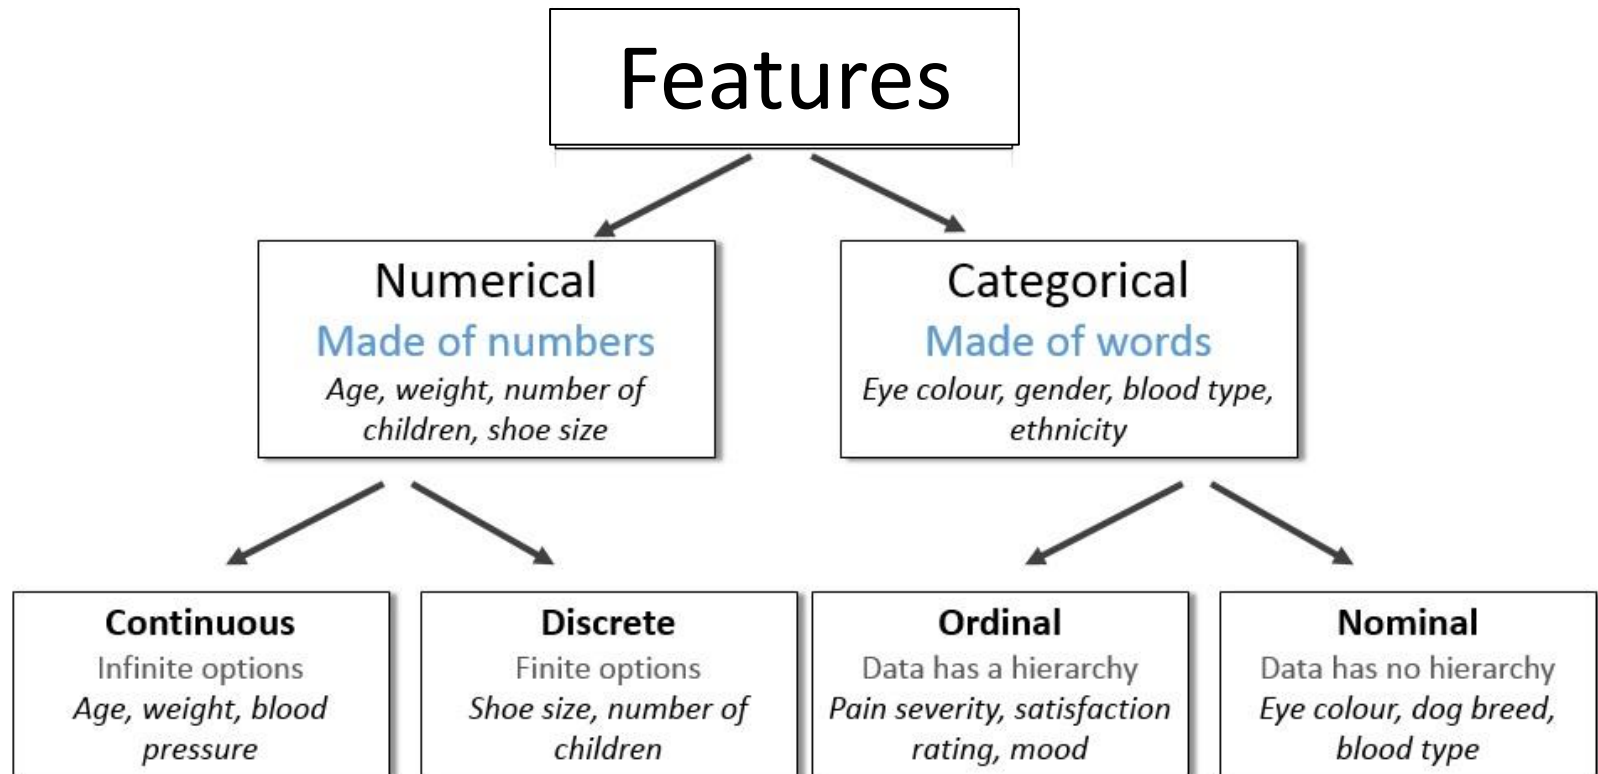
Reality :

- We have large amount of data but not *good enough*
- How to transform your data into learning compactable ?

What is a Feature in ML

- A **feature** is a measurable property of the object you're trying to analyze. In datasets, features appear as columns
- **Feature engineering** is the process of transforming raw data into features that better represent the underlying problem to the predictive models, resulting in improved model accuracy on unseen data.
- **Feature engineering** turn your inputs into things the algorithm can understand

Types of Features



Features from observation

An Apple



How to describe this picture?

Features from observation

More Fruits

- **Method I:** Use size of picture



(640, 580)



(640, 580)

- **Method II:** Use RGB average



(219, 156, 140)



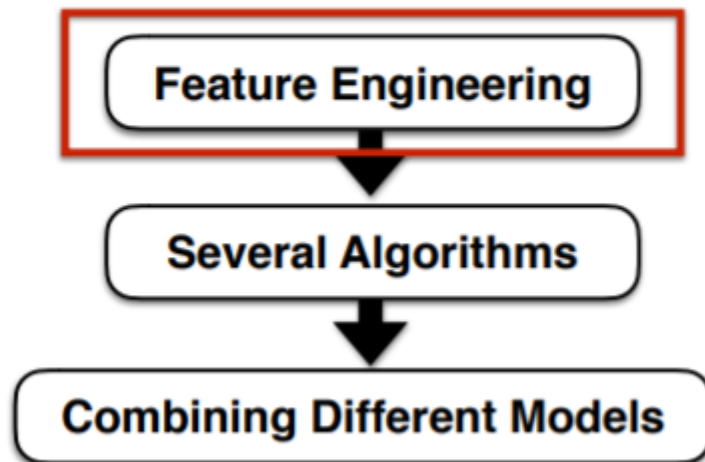
(243, 194, 113)




(216, 156, 155)

Feature Engineering


Pipeline



Feature Engineering



Features in iris data set



Source : Kaggle
<https://www.kaggle.com/uciml/iris>

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

first 5 rows of the Iris Dataset

Iris dataset in scikit-learn

- **scikit-learn** – machine learning in python
- Simple and efficient tools for data mining and data analysis
- several machine learning algorithms

```
from sklearn import datasets
```

```
# import some data to play with
```

```
iris = datasets.load_iris()
```

```
print(iris.data)
```


```
print(iris.target)
```

```
[[5.1 3.5 1.4 0.2]
 [4.9 3.  1.4 0.2]
 [4.7 3.2 1.3 0.2]
 [4.6 3.1 1.5 0.2]
 [5.  3.6 1.4 0.2]
 [5.4 3.9 1.7 0.4]
 [4.6 3.4 1.4 0.3]
 [5.  3.4 1.5 0.2]
 [4.4 2.9 1.4 0.2]
 [4.9 3.1 1.5 0.1]]
[0 0 0 0 0 0 0 0 0 0]
```

Dimensionality Reduction

Dimensionality reduction – Reducing the number of random variables to consider.

- Simple example
 - 3-D data



Dimensionality Reduction

- ‘**Dimensionality**’ - simply refers to the number of features (i.e. input variables) in your dataset.
- When the number of features is very large relative to the number of observations in your dataset, certain algorithms struggle to train effective models.
- This is called the “Curse of Dimensionality,”
- It’s especially relevant for clustering algorithms that rely on distance calculations.


Components of Dimensionality Reduction

Feature selection: you select a subset of the original feature set.



Feature extraction: you build a new set of features from the original feature set.

Methods of Dimensionality Reduction

- Principal Component Analysis (PCA)
- Linear Discriminant Analysis (LDA)



- Reduce dimensionality, preserve as much class discriminatory information as possible



Feature Reduction Iris Dataset

```
import matplotlib.pyplot as plt

from sklearn import datasets
from sklearn.decomposition import PCA
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

iris = datasets.load_iris()

X = iris.data
y = iris.target
target_names = iris.target_names

pca = PCA(n_components=2)
X_r = pca.fit(X).transform(X)
print(X_r)

lda = LDA(n_components=2)
X_r2 = lda.fit(X, y).transform(X)
print(X_r2)
```


Feature Reduction Iris Dataset

PCA OUTPUT

```
[[ -2.68412563  0.31939725] [-2.71414169 -0.17700123] [-2.88899057 -0.14494943] [-2.74534286 -  
0.31829898] [-2.72871654  0.32675451] [-2.28085963  0.74133045] [-2.82053775 -0.08946138] [-  
2.62614497  0.16338496] [-2.88638273 -0.57831175] -----  
-----  
-----  
-----[ 2.15943764 -0.21727758] [ 1.44416124 -0.14341341] [ 1.78129481 -0.49990168] [  
3.07649993  0.68808568] [ 2.14424331  0.1400642 ] [ 1.90509815  0.04930053] [ 1.16932634 -  
0.16499026] [ 2.10761114  0.37228787] [ 2.31415471  0.18365128] [ 1.9222678  0.40920347] [  
1.41523588 -0.57491635] [ 2.56301338  0.2778626 ] [ 2.41874618  0.3047982 ] [ 1.94410979  0.1875323  
] [ 1.52716661 -0.37531698] [ 1.76434572  0.07885885] [ 1.90094161  0.11662796] [ 1.39018886 -  
0.28266094]]
```

LDA OUTPUT

```
[[ -8.06179978e+00  3.00420621e-01] [-7.12868772e+00 -7.86660426e-01] [-7.48982797e+00 -  
2.65384488e-01] [-6.81320057e+00 -6.70631068e-01] [-8.13230933e+00  5.14462530e-01] [-  
7.70194674e+00  1.46172097e+00] [-7.21261762e+00  3.55836209e-01] -----  
-----  
-----  
-----[ 9.97610366e-01 -4.90530602e-01] [ 3.83525931e+00 -1.40595806e+00] [  
2.25741249e+00 -1.42679423e+00] [ 1.25571326e+00 -5.46424197e-01] [ 1.43755762e+00 -1.34424979e-  
01] [ 2.45906137e+00 -9.35277280e-01] [ 3.51848495e+00  1.60588866e-01] [ 2.58979871e+00 -  
1.74611728e-01] [-3.07487884e-01 -1.31887146e+00] [ 4.96774090e+00  8.21140550e-01] [  
5.88614539e+00  2.34509051e+00] [ 4.68315426e+00  3.32033811e-01]]
```

Preprocessing

Scaling

- The `MinMaxScaler` is the probably the most famous scaling algorithm, and follows the following formula for each feature:

$$\frac{x_i - \min(x)}{\max(x) - \min(x)}$$

- Shrinks the range such that the range is now between 0 and 1 (or -1 to 1 if there are negative values).

```
from sklearn.preprocessing import MinMaxScaler

data = [[-1, 2], [-0.5, 6], [0, 10], [1, 18]]
scaler = MinMaxScaler()
print(scaler.fit(data))
#MinMaxScaler(copy=True, feature_range=(0, 1))
print(scaler.transform(data))
#[[0.  0. ]
# [0.25 0.25]
# [0.5  0.5 ]
# [1.   1.   ]]
print(scaler.transform([[2, 2]]))
#[[1.5 0.  ]]
```

Preprocessing

Label Encoding

- Used to transform non-numerical labels that is categorical values to numerical labels.

```
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
le.fit(["paris", "paris", "tokyo", "amsterdam"])
enlabels=le.transform(["paris","tokyo", "tokyo", "paris"])
print(enlabels)
#[1 2 2 1]
print(le.inverse_transform(enlabels))
#['paris' 'tokyo' 'tokyo' 'paris']
```

Thank You

